


Background

- The use of **mathematical models** to create **forecasts** on the possible paths of **epidemics** and **pandemics** in almost real-time to inform public health interventions has gained more attention in **the past ten years**
- For example, we can point out **the US CDC Flu Sight Challenge**, **the DARPA Chikungunya Challenge**, **the Dengue Forecasting Challenge**, and **the Ebola Forecasting Challenge** along with epidemic and pandemic emergencies including **the 2014–16 West African Ebola epidemic**, **the 2018–19 DRC Ebola epidemic**, and **the ongoing COVID-19 pandemic**.



Deaths are gradually rising again

Number of daily reported coronavirus deaths in the US



Source: COVID Tracking Project

Frequentist and Bayesian analysis

- In the **Frequentist method (the QuantDiffForecast)**, our treatment of **parameter estimation** assumes that θ is an unknown but non-random quantity. It is some fixed parameter describing the true distribution of data, and our goal was to determine this parameter.
- **The Bayesian paradigm** naturally incorporates our prior belief about the unknown parameter θ , and updates this belief based on observed data.
- In Bayesian analysis, before data is observed, **the unknown parameter** is modeled as a **random variable** θ having a probability distribution Θ , called **the prior distribution**. This distribution represents our prior belief about the value of this parameter.
- The conditional distribution of Θ given the observed data is called **the posterior distribution**. It represents our knowledge about the parameter Θ after having observed the data.

Bayes' rule

- Inference reverse-engineers process the data and aim to estimate parameter values given the observations. In a Bayesian framework, the set of plausible parameter values conditional on the data is characterized by the posterior distribution. The posterior distribution combines information from the data and prior knowledge and the result is obtained via Bayes' rule.
- So, based on the Bayes' rule, we have

$$p(\theta|X) \propto p(\theta)p(X|\theta).$$

Or, we can simply say

$$\text{Posterior density} \propto \text{Prior density} \times \text{Likelihood}.$$

Algorithms

MCMC and HMC

- **Markov chain Monte Carlo (MCMC)** methods comprise a class of algorithms for **sampling** from a probability distribution. By constructing a **Markov chain**, one can obtain a sample of **the desired distribution** by recording states from the chain. The more **steps** that are included, **the more closely** the distribution of the sample matches the actual desired distribution.
- **The Hamiltonian Monte Carlo (HMC)** algorithm (originally known as hybrid Monte Carlo) is a **Markov chain Monte Carlo** method for obtaining a sequence of random samples that **converge** to being distributed according to a target probability distribution for which direct sampling is difficult. This sequence can be used to estimate integrals concerning the target distribution (expected values).

SEIR Model


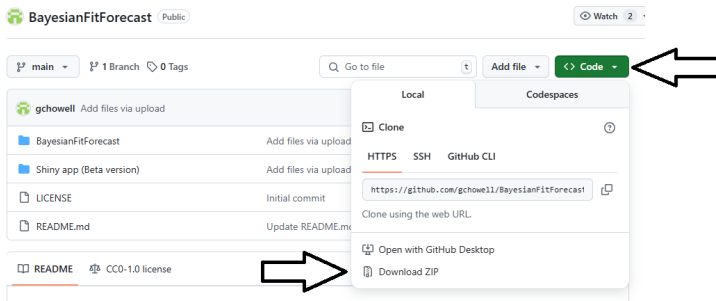


Figure: Compartmental diagram of the SEIR model with underreporting. Circles show the epidemiological compartments: susceptible (S), exposed (E), infectious (I), and recovered (R). Solid arrows indicate the transitions between compartments. The dashed arrow indicates the source of the observations, which are the newly reported infected individuals.

Downloading the Toolbox

- Google BayesianFitForecast GitHub.
- Click on the link to the GitHub repository.
- On the repository page, click the **Code** button.
- Select **Download ZIP** from the dropdown.



- A set of small navigation icons typically found in Beamer presentations, including symbols for back, forward, search, and other slide controls.

The `prior_sltn.R` Function

Purpose


Generates **prior predictive simulations** by sampling from user-defined prior distributions.

Workflow


- 1 Draw samples from prior distributions
- 2 Solve the ODE system using each sample
- 3 Construct an ensemble of trajectories to reflect **prior uncertainty**
- 4 To use this function, you must define the ODE model—including parameters and state variables, and specify the prior distributions for the parameters within the options file. The structure and setup of this file will be explained in detail later.

Prior Predictive Simulation

- Tight prior distributions yield low variability in trajectories
- Simulated trajectories represent our **initial beliefs** before observing data



(a) Incidence Cases



(b) All State Variables

Goal of the BayesianFitForecast Toolbox

Main Objective

Estimate the **posterior distribution** of model parameters that best explain the observed data.

What's Next

To achieve this, we must first fully explain the structure and role of the **options file**.

Next: Posterior Distribution

- We now want to **combine prior knowledge with observed data**
- Goal: obtain **posterior distributions** of parameters
- This requires the **options file** configuration

Reviewing all settings in one option file

Calibration Periods

Definition: The calibration periods are the specific time intervals used to fit the model to the data.

Input: Calibration periods should be entered as an array of numbers, where each number is selected from 1 to the total number of data points.

Important Note: Even if you want to consider only one calibration period, it should still be entered as an array with a single element.

Examples:

- **Multiple Calibration Periods:**

```
calibrationperiods <- c(10, 15, 20, 25, 30)
```

This indicates interest in fitting the model at several calibration periods: 10, 15, 20, 25, 30. Calibration period 10 means that we are fitting the first ten data points to the model.

- **Single Calibration Period:**

```
calibrationperiods <- c(10)
```

Forecasting Horizon

Definition: The "**forecastinghorizon**" defines how far into the future predictions are made after the model is fit. It impacts the accuracy and uncertainty of predictions.

Flexibility: By choosing a **large forecasting horizon initially**, you gain the flexibility to generate metrics for a smaller time frame later.

Adjusting the Forecasting Horizon:

- **Example:** If you set

```
forecastinghorizon <- 10
```

when running "**run_MCMC**", you can later generate results for any horizon smaller equal than 10, such as $1, 2, \dots, 10$. This allows you to focus on a specific part of the forecast while obtaining updated metrics. The reason to mention this point is that our toolbox is consisting of running two codes, and the first code takes much longer. So, having a large number for horizon makes our life easier for any smaller horizon later.

Adjusting the Forecasting Horizon

- **Adjusting Horizon in Post-Processing:** To zoom into a smaller segment of the forecast, change the

`forecastinghorizon`

in the options file **after** running "`run_MCMC.R`" but **before** executing "`run_analyzeResults.R`".

- **Note:** Results may differ if you initially set

`forecastinghorizon <- 5`

However, if the chains in a larger one are **convergent**, this difference is **tiny**.

- **Important:** The `forecastinghorizon` must always be a positive integer. Zero `forecastinghorizon` means fitting.

Model Name

Definition: The `model_name` parameter allows the user to assign a name to the project, helping avoid confusion when extracting the results.

Usage: Assigning distinct names to projects, especially when settings are similar, prevents overwriting results in the same folder. This is particularly useful when running models with different priors or parameters.

Examples:

```
model_name <- "Bayesian-normalpriors"  
model_name <- "Bayesian-niter=1000"
```

State Variables in the SEIR Model

Definition: The state variables in the SEIR model represent **different compartments** of the population and are defined as follows:

- "S": Susceptible
- "E": Exposed
- "I": Infected
- "R": Recovered
- "C": Cumulative number of infected individuals

Variables Definition: The variables should be listed **in the same order as their derivatives** appear in the system of ODEs.

Example:

```
vars <- c("S", "E", "I", "R", "C")
```

Model Parameters

Definition: The **parameters** in the SEIR model represent values that affect the dynamics of the system. They can be either **constants** or **values to be estimated** during model fitting.

Parameters:

- "beta": Transmission rate
- "gamma": Recovery rate
- "kappa": Incubation rate
- "rho": Reporting proportion
- "N": Total population size

Parameters Definition: All **parameters** must be included in the **vector params**, and each parameter should be defined as a string.

Example:

```
params <- c("beta", "gamma", "kappa", "rho", "N")
```


Defining the ODE System

Overview: The "ode_system" defines the system of ODEs by specifying the relationships between parameters, variables, and their derivatives with respect to time.

Notation:

- **"params*i*"** refers to the *i*-th parameter in the params list.
- **"vars*i*"** refers to the *i*-th variable in the vars list.
- **"diff vars*i*"** refers to the derivative for the *i*-th variable.

Formatting Rules:

- The first line of the `ode_system` string should be left empty.
- The last line of the string must end with the final equation, followed by the closing quotation mark.

Defining the ODE System

Example: Assume the following definitions for variables and parameters:

```
vars <- c("S", "E", "I", "R", "C")
```

```
params <- c("beta", "gamma", "kappa", "rho", "N")
```

Then, the ODE system can be defined as:

```
ode_system <- '
diff_var1 = -params1 * vars3 * vars1 / params5
diff_var2 = params1 * vars3 * vars1 / params5 - params3 * vars2
diff_var3 = params3 * vars2 - params2 * vars3
diff_var4 = params2 * vars3
diff_var5 = params4 * params3 * vars2'
```


Composite Expressions

Definition: Sometimes, users need to estimate derived parameters, such as the **basic reproduction number** or **recovery time**. To facilitate this, users should include these parameters in the `composite_expressions` list. Each entry should pair the parameter name on the left with its corresponding formula on the right, expressed as a string and a function of the model's parameters.

Example: In the SEIR model, the parameters might be defined as follows:

```
composite_expressions <- list(  
  R0 = "beta/gamma",  
  recovery_time = "1/gamma"  
)
```

fitting_index and fitting_diff

Definition: These parameters map the data to model variables. For example, in the SEIR model, if data represents the infectious individuals, set `fitting_index` to `c(3)`, where I is the 3rd variable in `vars`. For the cumulative cases in the SEIR model, set `fitting_index` to `c(5)` for C .

Derivatives: To fit data to a variable's derivative, set `fitting_diff` to `c(1)`. For instance, `fitting_index = c(5)` and `fitting_diff = c(1)` in SEIR fits data to $\frac{dC}{dt}$ (incidence cases), while `fitting_index = c(3)` and `fitting_diff = c(1)` fits to $\frac{dI}{dt}$ (newly infectious).

Multiple datasets: For multiple time-series, use arrays. For example, `fitting_index = c(3,5)` and `fitting_diff = c(0,1)` fit the data to I and $\frac{dC}{dt}$.

errstrc

Definition: The `errstrc` parameter is an integer that specifies the **error structure** used when fitting the model to data, which describes the variability or noise.

Options:

- 1 **Negative binomial**: accounts for overdispersion in count data.
- 2 **Normal**: assumes constant variance.
- 3 **Poisson**: assumes variance equals the mean.

Selection: Choose the error structure by setting `errstrc` to the corresponding number.

Additional Parameters:

- Negative binomial requires estimating dispersion ϕ .
- Normal requires estimating standard deviation σ .

cadfilename1 and caddisease

cadfilename1: This is a string representing the name of the Excel file (without the extension `.xlsx`). The Excel file should have at least two columns:

- **First column:** Sequential time points (e.g., 0, 1, 2, ...) representing days, weeks, or years.
- **Second column:** Temporal incidence data related to the state variables or their derivatives. The column name should start with **cases1**, followed by **cases2**, **cases3**, etc., for additional columns.

caddisease: This string defines the **name of the disease** or process. It is used to generate a unique name for the results folder, preventing overwriting results from different diseases.

series_cases and datatype

series_cases: This string parameter defines the time series label(s) for the y-axis. For example:

- If set to "Cases", the y-axis label will be "Cases" for a single time series.
- If set to "infected", "recovered", there are two time series, and the y-axis labels will be "infected" for cases1 and "recovered" for cases2.

datatype: This parameter specifies the x-axis label. For instance, setting "datatype" to "Days" labels the x-axis as "Days".

Priors - Example

Example: Given:

- `params <- c("beta", "gamma", "kappa", "rho", "N")`

Prior beliefs:

- `beta`: 2 (constant)
- `gamma, kappa`: `"normal(1,2)T[0,]"` (truncated normal)
- `rho`: `"uniform(0,1)"` (uniform)
- `N`: 10000 (constant)

The prior distribution setup:

- `params1_prior <- 2`
- `params2_prior <- "normal(1,2)T[0,]"`
- `params3_prior <- "normal(1,2)T[0,]"`
- `params4_prior <- "uniform(0,1)"`
- `params5_prior <- 10000`

- A set of small navigation icons typically found in Beamer presentations, including symbols for back, forward, search, and other slide controls.

Lower and Upper Bounds - Example

Example: Given:

- `params <- c("beta", "gamma", "kappa", "rho", "N")`

Parameter bounds:

- `beta`, `gamma`, `kappa`: Lower bounds 0, no upper bounds.
- `rho`: Lower bound 0, upper bound 1.
- `N`: No lower or upper bounds.

Bound definitions:

- `params1_LB <- 0, params1_UB <- NA`
- `params2_LB <- 0, params2_UB <- NA`
- `params3_LB <- 0, params3_UB <- NA`
- `params4_LB <- 0, params4_UB <- 1`
- `params5_LB <- NA, params5_UB <- NA`

- ϕ for the negative binomial.
- σ for the normal distribution.

- `normalerror1_prior <- "cauchy(0, 2.5)T[0,]"`
(Cauchy distribution with location 0 and scale 2.5).

Prior Distribution for Additional Parameters in Error Structure

Multiple Datasets: For three time series, define:

- `normalerror1_prior <- "cauchy(0, 2.5)T[0,]"`
- `normalerror2_prior <- "cauchy(0, 2.5)T[0,]"`
- `normalerror3_prior <- "cauchy(0, 2.5)T[0,]"`

Error Structure Setting:

- If `errstrc` = 1, define only `negbinerror_prior`.
- If `errstrc` = 2, define only `normalerror_prior`.

Initial Conditions (Ic)

Definition: The vector `Ic` represents the initial conditions for a Bayesian method in Stan, corresponding to initial state variable values. The order must match `vars`.

Example (SEIR Model):

$$S(0) = N - i_0, \quad E(0) = 0, \quad I(0) = i_0, \quad R(0) = 0, \quad C(0) = i_0$$

where i_0 is the first observed data point and N is the population size.

If $i_0 = 1$:

$$lc = c(params5_prior - 1, 0, 1, 0, 1)$$

Alternatively, specify N directly:

$$l_c = c(10000 - 1, 0, 1, 0, 1)$$


```
vars.init
```

Purpose: Sometimes, the **initial conditions** (e.g., initial number of infected or exposed individuals) are unknown and need to be estimated. This can be done by setting **vars.init = 0**. If the values are known, set **vars.init = 1**.

When vars.init = 0:

- **params**: Add the new parameter (e.g., i_0) to params.
- **paramsfix**: Update to include the new parameter, with its entry set to 0 (for estimation).
- **priors**: Define the prior distribution for the new parameter.
- **Ic**: Specify initial conditions. Example for SEIR:

$$I_C = (N - i_0, 0, i_0, 0, i_0)$$

When `vars.init = 1`: No changes are needed.

◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡

Process

**What is the process after
configuring the option file?**

run_MCMC

Purpose: The file `run_MCMC` is designed to fit the model to the data using user-defined parameters.

Configuration:

- The user must update the options in the **options** file as needed.
- The user must also ensure that **line 5** in the `run_MCMC` file matches the name of the configured options file.

Importance: Updating the options and ensuring correct file references is crucial for running the MCMC fitting process smoothly.

Saved Data

Output: After running `run_MCMC`, the generated samples will be saved as an `.Rdata` file in the toolbox folder.

File Naming: For each period in `calibrationperiods`, a corresponding `.Rdata` file is generated. The filename is based on:

- model_name
- calibrationperiod
- forecastinghorizon
- errstrc
- caddisease

Recommendation: If two option files share identical settings, the second run will **overwrite** the first file. It is recommended to modify **model_name** to avoid this.

run_analyzeResults

Purpose: By running `run_analyzeResults`, users can generate and analyze results based on the previously configured options.

Key Points: Ensure the `option file` remains the same as in `run_MCMC` (except for the forecasting horizon). Modify line 10 to specify the option file name.

File Naming: Generated files will follow the format:

```
model_name + caddisease + errstrc +  
calibrationperiod + forecastinghorizon
```

Convergence and Forecast Files

Convergence File: Includes the convergence parameter **Rhat**. Use results with **Rhat** < 1.1 . To achieve this, increase **niter** or **num_chain**.

Example format:

`convergence-Bayesian-weak-sanfrancisco-normal-cal-10-fcst-10`

Forecast File: Includes columns for Date, Data, Median, Lower/Upper Bounds of 95% Pls. Example format:

`forecast-Bayesian-weak-sanfrancisco-normal-cal-10-fcst-10`

Parameters and Performance Metrics

Parameters File: Provides estimates with median, mean, lower/upper bounds of 95% CIs. For normal error structure: σ , For negative binomial: ϕ .

Example format:

`parameters-Bayesian-weak-sanfrancisco-normal-cal-10-fcst-10`

Performance Metrics File: Includes MAE, MSE, WIS, and Coverage for calibration/forecasting periods. Example format: `performance`

`metrics-Bayesian-weak-sanfrancisco-normal-cal-10-fcst-10`

Generated PDF Files

In addition to Excel files, several PDF files will be generated:

- **Histograms:** Save histograms of parameters, including σ or ϕ . Example: `"beta-histogram-Bayesian-weak-\sanfrancisco-normal-cal-10-fcst-10"`.
- **Forecast:** Shows median fit and 95% PI margins. Example: `"Forecast-Bayesian-weak-sanfrancisco-normal-\cal-10-fcst-10"`.
- **Trace plot:** Visualizes parameter values over MCMC iterations. Example: `"traceplot-Bayesian-weak-sanfrancisco-normal-cal-10-fcst-10"`.

Key parameters in the new file

```
calibrationperiods <- c(17)
forecastinghorizon <- 10
vars <- c("S", "E", "I", "R", "C")
ode_system <- '
  diff_var1 = -params1 * vars3 * vars1 / params5
  diff_var2 = params1 * vars3 * vars1 / params5 - params3 * vars2
  diff_var3 = params3 * vars2 - params2 * vars3
  diff_var4 = params2 * vars3
  diff_var5 = params4 * params3 * vars2'
paramsfix <- c(0,1,1,1,1)
```

Settings and execution

Important Settings

Ensure the following settings are correct for fitting and error structure:

- `fitting_index = 5`
- `fitting_diff = 1`
- `errstrc = 3`

Execution

To run the MCMC algorithm:

- Modify line 5 in `run_MCMC` to use the new options file:
"options_SEIR_sanfrancisco_Ex1.R".
- Execute the script in RStudio and analyze results using `run_analyzeResults.R`.

Fitting using Poisson error structure






Figure: Bayesian fitting of the SEIR model to the 1918 influenza dataset from San Francisco was performed using a Poisson error structure, calibrated over 17 days. The model captures the dynamics of newly infected individuals, represented by $\frac{dC}{dt}$. The parameters $\kappa = \frac{1}{1.9}$, $\gamma = \frac{1}{1.4}$, and $\rho = 1$ are held constant, while the prior distribution for β is uniform(0,10). All parameters were constrained to have non-negative values, with a total population size fixed at 550,000. Initial conditions were set as (549996, 0, 4, 0, 4), based on the first recorded data point. The MCMC algorithm was executed with 1,000 iterations across two independent chains to ensure robust convergence.

Histogram of posterior distributions: Poisson error structure



Posterior inference and convergence analysis: Poisson error structure

Table: Posterior inference and convergence analysis for the example applying the SEIR model to the San Francisco 1918 influenza dataset with a Poisson error structure.

Calibration	Parameter	Mean	Median	CI_95	N_eff	Rhat
17	β	0.71	0.71	(0.7 , 0.73)	316.35	1

Performance metrics: Poisson error structure

Table: The performance metrics resulting from Bayesian fitting of the SEIR model to the 1918 influenza dataset from San Francisco, using a Poisson error structure and 17 days of calibration data.

	Calibration	Forecasting
MAE	5.24	NA
MSE	45.35	NA
WIS	3.33	NA
Coverage of 95% CI	88.24	NA

Example 2

This example compares the performance metrics of the SEIR model using the SF 1918 flu dataset with two error structures: negative binomial and Poisson. The top panel shows the model fit, while the bottom panels display histograms of the estimated parameters: β , ϕ , and R_0 .

Additionally, the performance metrics and convergence analysis are presented in tables for a detailed comparison of the two error structures. The option file for this example can be found under the name "options_SEIR_sanfrancisco_Ex2".

Fitting using negative binomial error structure







Figure: The Bayesian fitting of the SEIR model to the first 17 days of the 1918 influenza pandemic in San Francisco was performed using a negative binomial error structure. The model is fitted to the newly infected individuals, represented by $\frac{dC}{dt}$. The parameters $\kappa = \frac{1}{1.9}$, $\gamma = \frac{1}{1.4}$, and $\rho = 1$ are constants, while the prior distribution for β is $\text{uniform}(0,10)$ and for ϕ is $\text{exponential}(5)$, with all parameters having a lower bound of zero. The population size is 550,000, with an initial condition of (549996, 0, 4, 0, 4) based on the first recorded data point. The MCMC algorithm was run with 1,000 iterations across two chains.

Parameter histograms: negative binomial error structure



Performance metrics: negative binomial error structure

Table: Performance metrics for the SEIR model with a negative binomial error structure, based on the San Francisco 1918 influenza dataset, calibrated over 17 days.

	Calibration	Forecasting
MAE	5.76	NA
MSE	68.24	NA
WIS	3.57	NA
Coverage of 95% PI	100.00	NA

Convergence: negative binomial error structure

Table: Convergence and other statistics for the SEIR model applied to the San Francisco 1918 influenza dataset with a negative binomial error structure.

Parameter	Mean	Median	CI_95	N_eff	Rhat
β	0.73	0.73	(0.7 , 0.78)	356.83	1
ϕ	20.76	12.65	(3.65 , 92.49)	395.63	1

Comparison of performance metrics in calibrating


Table: Comparison of performance metrics between the Poisson and negative binomial error structures for the SEIR model, based on the SF 1918 influenza dataset.

Model	MAE	MSE	Coverage 95% PI	WIS
Bayesian, Poisson error	5.24	45.35	88.24	3.33
Bayesian, Negative binomial error	5.76	68.24	100	3.57


Example 3: Forecasting with Poisson vs. negative binomial

In this example, we present the best model for forecasting 10 days after calibrating 17 days of the SF 1918 flu dataset. During the calibration period, the Poisson error structure performs better in terms of MAE, MSE, and WIS metrics, while the negative binomial model excels in 95% PI coverage. However, the negative binomial model demonstrates superior performance across all metrics in the forecasting period. The forecast for both Poisson and negative binomial error structures is presented, along with a comparison of their performance metrics. The related option files for this example are provided as "options_SEIR_sanfrancisco_Ex3_Negbin" and "options_SEIR_sanfrancisco_Ex3_Poisson".

Negative binomial error structure



Poisson error structure



Comparison of performance metrics in forecasting

Table: A comparison of forecasting and model fit performance metrics between the Poisson and negative binomial error structures.

Error structure	MAE	MSE	Coverage 95% PI	WIS
Calibration performance				
Negative binomial	5.85	68.25	100.0	3.62
Poisson	5.00	41.71	88.24	3.28
Forecasting performance				
Negative binomial	92.10	12325.90	100.0	56.84
Poisson	149.75	29976.22	10.0	122.28



Example 4: Comparison of SEIR and exponential growth models

In this practice example, we compare the SEIR model to the exponential growth model using the SF 1918 flu data, with a negative binomial error structure. The models were calibrated for 17 days, followed by a forecast of 10 days.

The forecasting results, along with the histogram of the parameter r from the exponential growth model, are presented.

The results indicate that the SEIR model outperforms the exponential growth model in calibration performance across all metrics. However, the exponential growth model demonstrates superior performance in forecasting for all metrics.

Forecast using exponential growth model



Example 5: Forecasting the 1918 flu pandemic in San Francisco estimating the initial number of infected people

In this example, we utilize data from the 1918 influenza pandemic in San Francisco to forecast the number of infected individuals using a normal error structure. The model is calibrated over a 17-day period, followed by a 10-day forecast.

In this scenario, we estimate the parameter i_0 , which represents the initial number of infected individuals. The estimation reveals that i_0 is approximately 11.


The results demonstrate the model's effectiveness in capturing the dynamics of the outbreak and providing insights into the initial spread of the infection.

Forecast while estimating the initial number of infected people




Parameter histograms


Median: 0.57 Mean: 0.57 95% CI: 0.51 , 0.63



Median: 4.76 Mean: 4.88 95% CI: 3.46 , 7.18



Median: 11.05 Mean: 11.4 95% CI: 6.68 , 18



Example 6: Simulation of SEIR using multiple datasets

In this example, we simulate data using the forward solution of the SEIR model with the following parameters:

- $\beta = 0.5$
- $\gamma = 0.25$
- $\kappa = 1$
- $\rho = 1$
- $N = 100,000$

Normal noise with a standard deviation of 5 is added to the simulation.

The simulated data includes:

- Day (Column 1)
- Number of infectious individuals I (Column 2)
- Derivative of recovered individuals $\frac{dR}{dt}$ (Column 3)
- Derivative of cumulative cases $\frac{dC}{dt}$ (Column 4)

Model Fitting Results and Analysis

Three modeling scenarios are considered:

- ① Fitting the SEIR model to all data
- ② Fitting the model to I and $\frac{dC}{dt}$
- ③ Fitting the model to only $\frac{dC}{dt}$




The calibration period spans 70 days, followed by a 30-day forecast. The fitted curves of the incidence, along with prediction bands for the three scenarios, highlight the following:

- Scenario 1, using all data, outperforms Scenario 2, which uses I and $\frac{dC}{dt}$.
- Scenario 2 performs better than Scenario 3, which relies solely on $\frac{dC}{dt}$.

Notably, while Scenario 3 performs better during calibration, it underperforms in forecasting. This suggests:

- Overfitting during calibration due to limited data.
- Scenario 1 provides more accurate parameter estimates, improving forecasting performance.

Forecast using three scenarios



Performance metrics using three scenarios

Table: Performance metrics for the rate of new cases, $\frac{dC}{dt}$, obtained by fitting the SEIR model to simulated data generated using the forward solution of the SEIR model.

Scenario	Period	MAE	MSE	Coverage 95% PI	WIS
Calibration Performance					
(i) $I, \frac{dR}{dt}, \frac{dC}{dt}$	30	3.06	15.59	96.67	1.99
(ii) $I, \frac{dC}{dt}$	30	3.22	16.02	93.33	2.02
(iii) $\frac{dC}{dt}$	30	3.04	14.08	96.67	1.90
Forecasting Performance					
(i) $I, \frac{dR}{dt}, \frac{dC}{dt}$	70	100.98	17216.70	87.14	65.78
(ii) $I, \frac{dC}{dt}$	70	131.93	35634.15	84.29	93.13
(iii) $\frac{dC}{dt}$	70	431.25	464089.54	60.0	287.49

Example 7: Time-Dependent SEIR Model

Goal: Model interventions that change transmission over time

Time-dependent transmission rate:

$$\beta(t) = \begin{cases} \beta_0, & \text{if } t < t_{\text{int}} \\ \beta_1 + (\beta_0 - \beta_1)e^{-q(t-t_{\text{int}})}, & \text{if } t \geq t_{\text{int}} \end{cases} \quad (1)$$

- β_0 : initial transmission rate
- β_1 : final transmission rate
- t_{int} : intervention time
- q : transition rate

Dataset Structure

File: `seir_simulated.xlsx`

- Column 1: Days
- Column 2: $\frac{dC}{dt}$ (new cases)

Options file: `options_seir_timedep_Ex7.R`

Task: Estimate β_0 , β_1 , q , t_{int} from data

Model Fit Results

